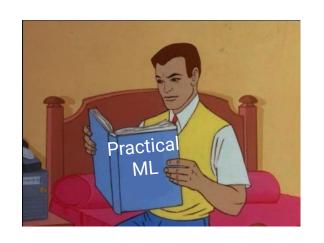
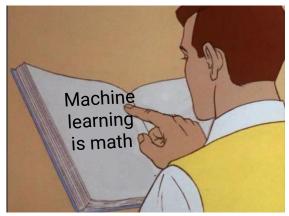
# Practical ML

Antonio Pitasi Samuele Sabella

2019 - Polo Fibonacci

# Before we get started, we need some theory







### Machine learning

- Practice: We define machine learning as a set of methods that can automatically detect patterns in data, and then use the uncovered patterns to predict future data, or to perform other kinds of decision making under uncertainty
  - Machine Learning: A Probabilistic Perspective, Kevin P. Murphy
- **Theory:** How does learning performance vary with the number of training examples presented? Which learning algorithms are most appropriate for various types of learning tasks?
  - Machine Learning, Tom Mitchell

# ML is not only artificial neural networks

- Lots of mathematical models
  - Hidden Markov models
  - Support Vector Machines
  - Decision trees
  - Boltzmann machines, Deep belief network, Deep Boltzmann
- Neural network models are many...
  - Shallow network, Deep neural network
  - CNN (Yolo, AlexNet, GoogLeNet)
  - Echo state network, Deep echo state network
  - o Rnn, LSTM, GRU

### Machine learning categories

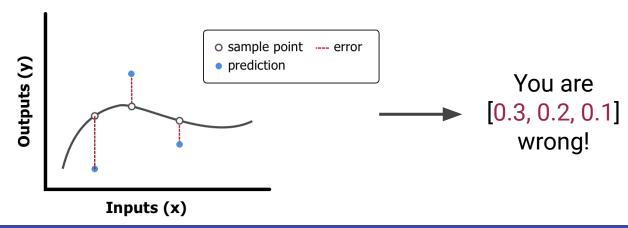
• **Supervised:** the goal is to learn a mapping from input to output given a dataset of labeled pairs  $\mathcal{D} = \{x_i, y_i\}_{i=1}^N$  called training set (e.g. Iris Data Set [2])

sepal length	sepal width	petal length	petal width	species
4.6	3.4	1.4	0.3	Iris-setosa
6.2	3.4	5.4	2.3	Iris-virginica

• **Unsupervised**: we have only a set of data points  $\mathcal{D} = \{x_i\}_{i=1}^N$  and the goal is to find interesting patterns in the data Example: young American males who buy diapers also have a predisposition to buy beer (original story [3])

#### How does it work?

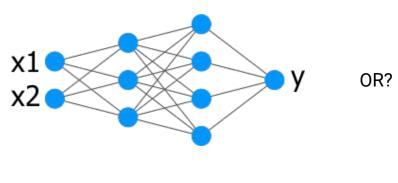
- Dataset of examples to learn from
- A model to learn from that data (e.g. neural net)
  - With some parameters to tune
  - With some hyperparameter to choose (neurons, layers, ...)
- Target function (loss) to minimize

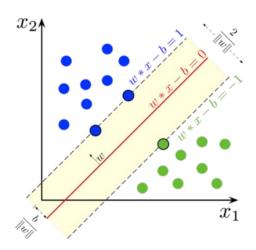




# What is usually done

- Validation phase: compare different models and configurations
  - Which model to choose

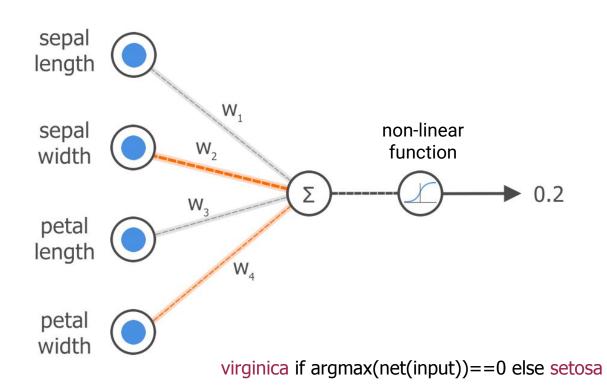


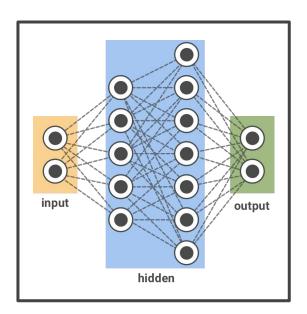


- Model hyper-parameters
- Test phase: loss, accuracy, recall, precision...
- We skip all this for seek of simplicity

Note: train/validation/test on different data

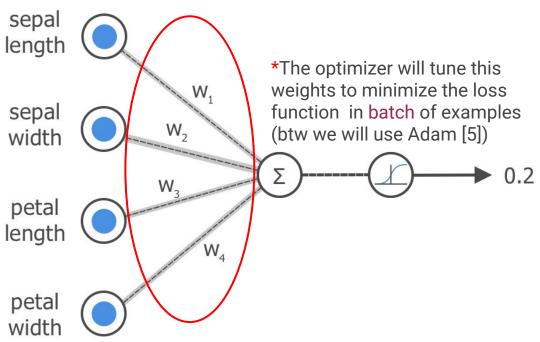
#### Models: feed-forward neural networks

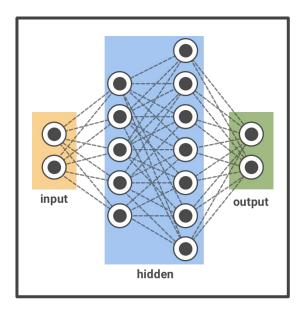




Stack neurons in layers

#### Models: feed-forward neural networks





Stack neurons in layers

virginica if argmax(net(input))==0 else setosa

#### A lot more stuff to know but for us...




UNDERSTANDING MACHINE LEARNING




import keras

# Practical ML








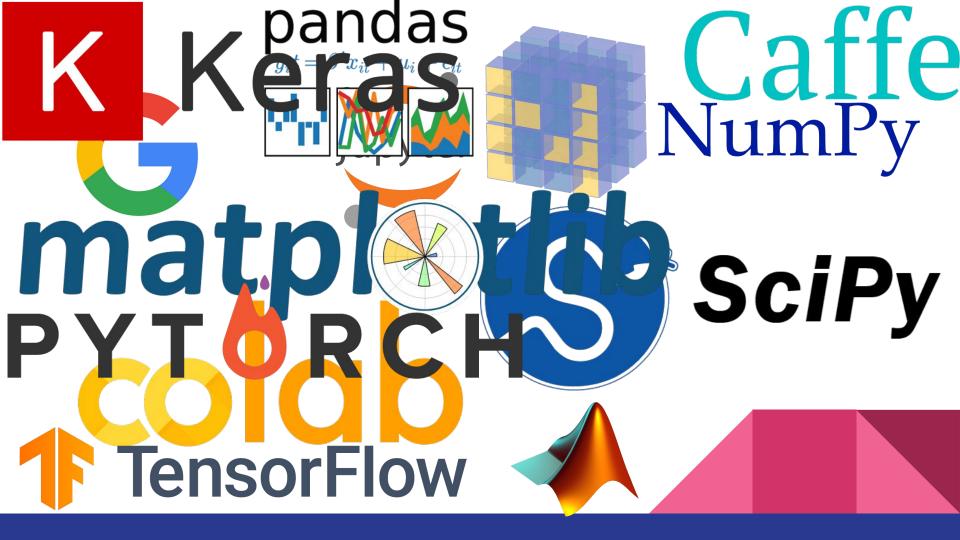


- Interactive
- Collaborative
- Python, R, Julia, Scala, ...





# coldb



# Keras

#### **Features:**

Easy to build a neural network

Easy to build a neural network wrong

# Keras

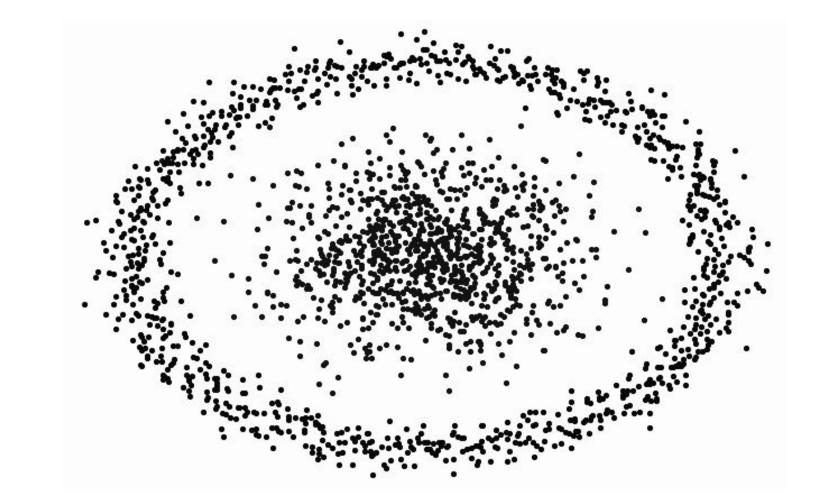
#### Keep an eye for:

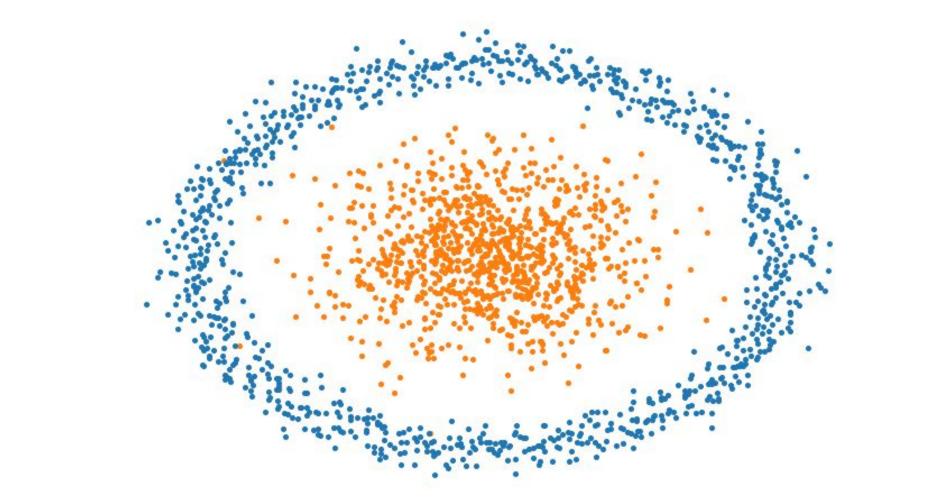
Accuracy

**Fitting** 

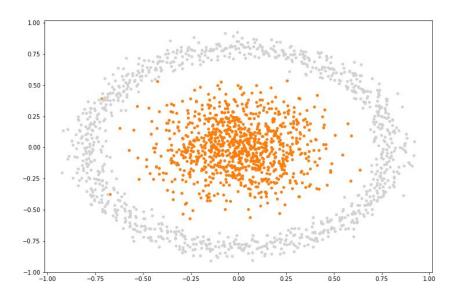
Performance

# Problem 1 **Points classification**



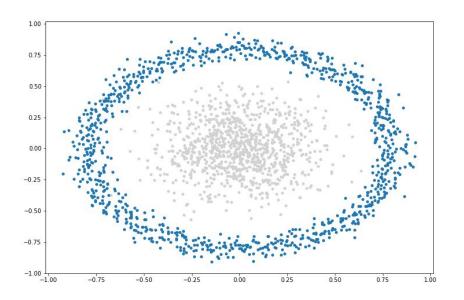


# Generating the dataset



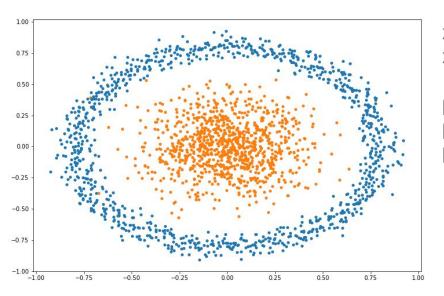
```
def make_inner(mu, sigma, num=1000):
    x = np.random.normal(mu, sigma, num)
    y = np.random.normal(mu, sigma, num)
    return x, y
```

#### Generating the dataset



```
def make_circle(mu, sigma, num=1000):
    r = np.random.normal(mu, sigma, num)
    phi = np.linspace(0,2.*np.pi, len(r))
    x = r * np.sin(phi)
    y = r * np.cos(phi)
    return x,y
```

# **Plotting**



```
x_inner, y_inner = make_inner(0, 0.2, n_points)
x_circle, y_circle = make_circle(0.8, 0.05, n_points)
```

```
plt.figure(figsize=(12,8))
plt.plot(x_inner, y_inner, 'o', markersize=4, c="C1")
plt.plot(x_circle, y_circle, 'o', markersize=4, c="C0")
```

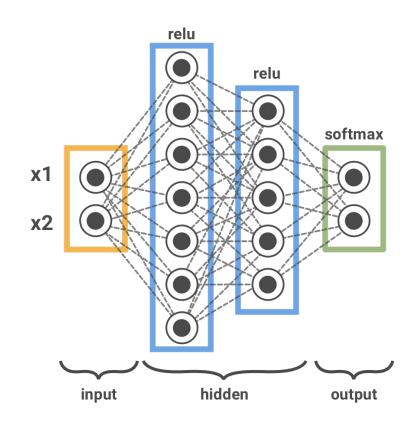
#### Our model

For non-linearity: rectifier linear unit

$$relu(x) = max(0, x)$$

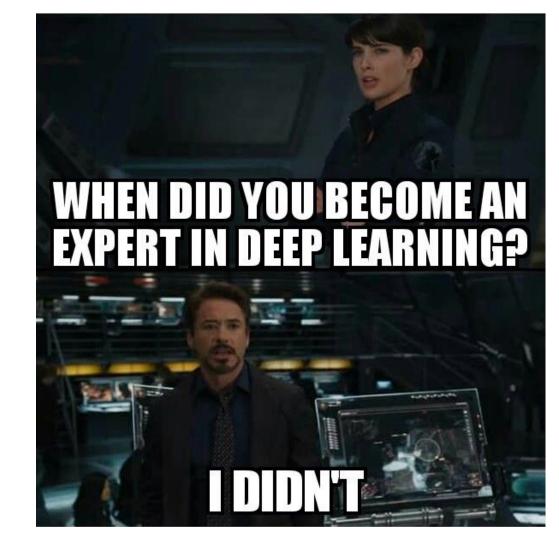
 We use a softmax function in the output layer to represent a probability distribution

$$softmax(x)_i = \frac{exp(x_i)}{\sum_j exp(x_j)}$$



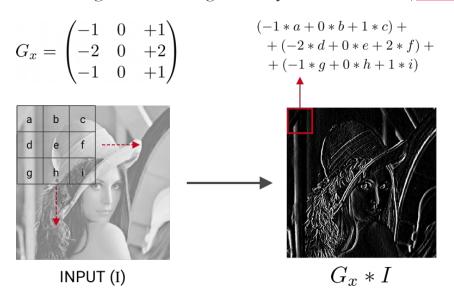
# Let's code!

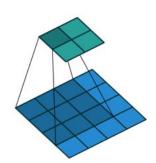
https://colab.research.google.com https://ml.anto.pt Me after training a neural network



# Back to theory - Convolving Lenna

• Given a function f, a convolution g with a kernel w is given by a very complex formula with a very simple meaning: "adding each element of the image to its local neighbors, weighted by the kernel" (wikipedia)

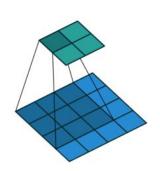




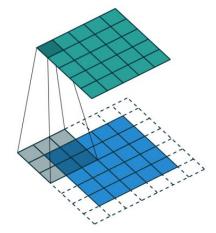
$$w * f(x,y) = \sum_{s=-a}^{a} \sum_{t=-b}^{b} w(s,t) f(x-s,y-t)$$

#### Convolution arithmetic

- Zero-padding: deal with borders pixels by adding zeros (preserves the size)
- Pooling: helps the network to become transformation invariant (translations, rotations...)



No padding, no strides

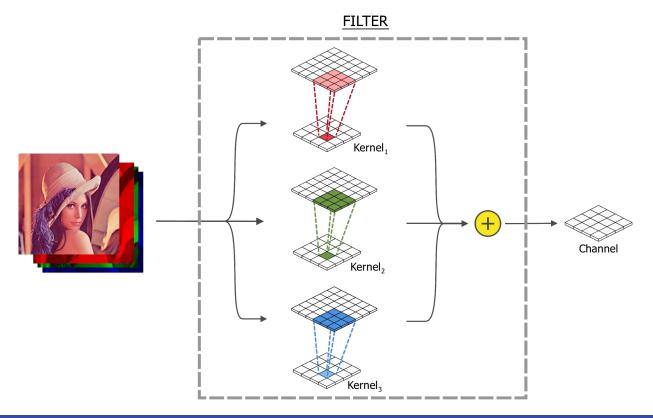


padding=same && no strides

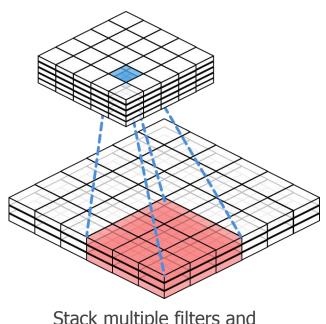
1	2	3	6		
1	2	3	3	2	6
4	5	4	1	5	4
5	0	2	3		

max pooling && 2x2 strides

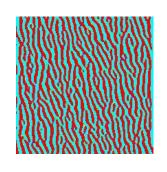
# Dealing with multiple input channels

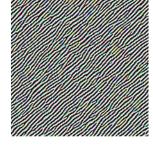


#### GoogLeNet on ImageNet - Feature visualization



Stack multiple filters and learn kernels dynamically (hierarchy of features)





feature visualization of the 1s conv. layer



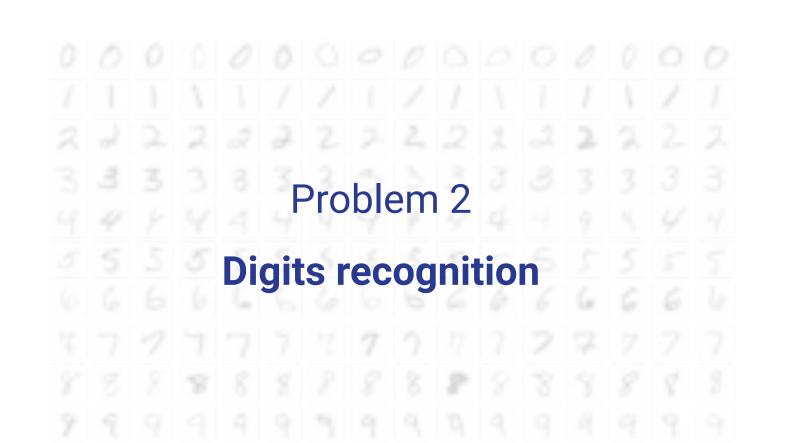








layer 4d



#### References

- [1] Pattern Recognition in a Bucket
- https://link.springer.com/chapter/10.1007/978-3-540-39432-7 63
- [2] Iris dataset: <a href="https://archive.ics.uci.edu/ml/datasets/iris">https://archive.ics.uci.edu/ml/datasets/iris</a>
- [3] Beer and diapers: <a href="http://www.dssresources.com/newsletters/66.php">http://www.dssresources.com/newsletters/66.php</a>
- [4] Multilayer feedforward networks are universal approximators:
- http://cognitivemedium.com/magic\_paper/assets/Hornik.pdf
- [5] Adam: A Method for Stochastic Optimization: <a href="https://arxiv.org/abs/1412.6980">https://arxiv.org/abs/1412.6980</a>
- [6] MNIST dataset: <a href="http://yann.lecun.com/exdb/mnist/">http://yann.lecun.com/exdb/mnist/</a>

#### References

[7] Bengio, Yoshua, Ian Goodfellow, and Aaron Courville. *Deep learning*. Vol. 1.

MIT press, 2017: <a href="http://www.deeplearningbook.org/">http://www.deeplearningbook.org/</a>

[8] Feature-visualization: <a href="https://distill.pub/2017/feature-visualization/">https://distill.pub/2017/feature-visualization/</a>

[9] Going deeper with convolutions: <a href="https://arxiv.org/pdf/1409.4842.pdf">https://arxiv.org/pdf/1409.4842.pdf</a>

[10] Imagenet: A large-scale hierarchical image database:

http://www.image-net.org/papers/imagenet\_cvpr09.pdf

[11] Culture, Communication, and an Information Age Madonna:

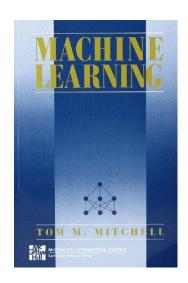
http://www.lenna.org/pcs\_mirror/may\_june01.pdf

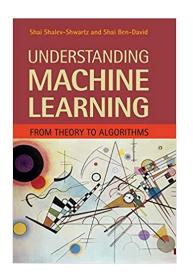
#### References

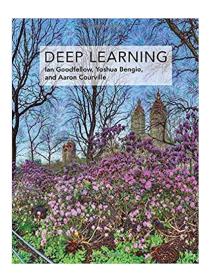
[12] Intuitively Understanding Convolutions for Deep Learning:

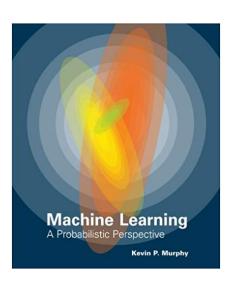
https://towardsdatascience.com/intuitively-understanding-convolutions-for-deep-learning-1f6f42faee1

#### **Books**







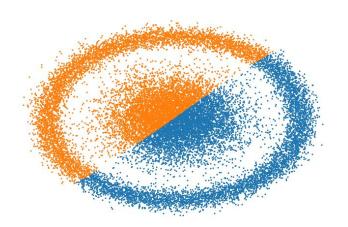


Difficulty

Antonio Pitasi Software Engineer, Nextworks https://anto.pt

Leave your feedback

Samuele Sabella <a href="https://github.com/samuelesabella">https://github.com/samuelesabella</a>





#### Antonio Pitasi

Software Engineer, Nextworks <a href="https://anto.pt">https://anto.pt</a>

# Samuele Sabella

https://github.com/samuelesabella

